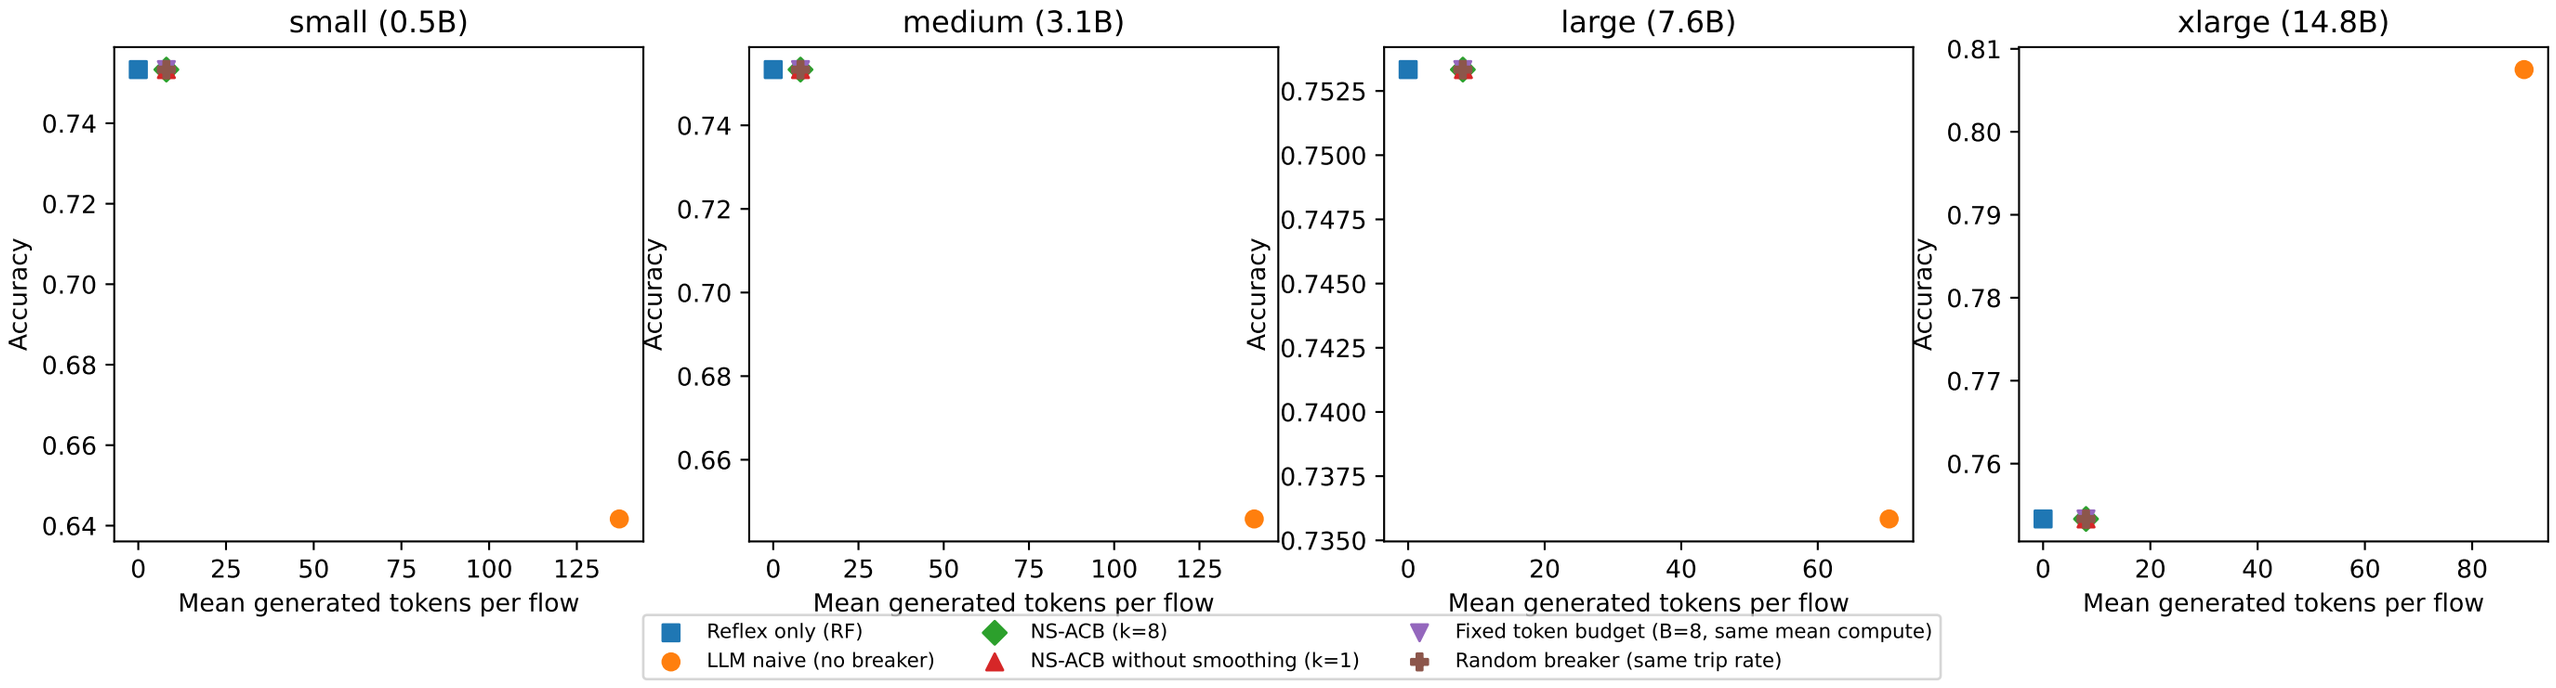
Accuracy vs. LLM compute (test flows, all datasets)
small (0.5B)
medium (3.1B)
large (7.6B)
xlarge (14.8B)
Accuracy
Mean generated tokens per flow
Reflex only (RF)
LLM naive (no breaker)
NS-ACB (k=8)
NS-ACB without smoothing (k=1)
Fixed token budget (B=8, same mean compute)
Random breaker (same trip rate)